

Scaling Hidden Markov Models

April 28, 2021

Latent Variables Models in NLP

- ▶ NLP benchmarks are dominated by fully observed models
 - ▶ Transformers
 - ▶ Previously, Recurrent Neural Networks
- ▶ We instead explore latent variable models

Latent Variable Models: Motivation

- ▶ LVMs posit a generative process involving unseen variables
- ▶ Maintain uncertainty over latent representations, rather than just output correlations
- ▶ Often improves interpretability and controllability
- ▶ Bottlenecked by the computational complexity of inference

Research Question

To what extent is the performance of tractable latent variable models limited by scale and choices in parameterization?

This work: Scale hidden Markov models (HMMs) on language modeling using techniques drawn from recent advances in neural networks

Language Modeling

How now, brown _____

- ▶ Given the words seen so far, predict the next word
- ▶ Language requires modeling long-range phenomena

Hidden Markov Models in NLP

- ▶ Simplest latent variable models for time series data
- ▶ Are thought to be very poor language models
- ▶ We show they are better than previously thought, once scaled

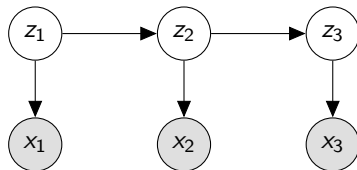
Background: Hidden Markov Models

Hidden Markov Models (HMMs)

- ▶ Classical models for unsupervised per-word tag induction
 - ▶ Part-of-speech induction
 - ▶ Word alignment for translation
- ▶ Admits tractable exact inference
 - ▶ Strong conditional independence assumptions
 - ▶ Finite set of discrete latent states

Hidden Markov Models (HMMs)

For times t , model states $z_t \in [Z]$, and tokens $x_t \in [X]$,



This yields the joint distribution

$$p(x, z) = \prod_t p(x_t | z_t) p(z_t | z_{t-1})$$

with

| | |
|---------------|---------------------|
| start state | $p(z_1),$ |
| transitions | $p(z_t z_{t-1}),$ |
| and emissions | $p(x_t z_t)$ |

represented as vectors and matrices

Inference

Given observed $x = (x_1, \dots, x_T)$ We wish to maximize

$$p(x) = \sum_{z_1} \cdots \sum_{z_T} p(x, z) = \alpha_1^\top \Lambda_2 \Lambda_3 \cdots \Lambda_T \mathbf{1},$$


where we have the

$$\begin{aligned} \text{start,} \quad & [\alpha_1]_{z_1} = p(x_1 \mid z_1)p(z_1), \\ \text{and transition operators,} \quad & [\Lambda_t]_{z_{t-1}, z_t} = p(x_t \mid z_t)p(z_t \mid z_{t-1}) \end{aligned}$$

The result of each matvec has the alphas of the forward algorithm, i.e. $\alpha_3 = \alpha_1 \Lambda_2 \Lambda_3$ has entries corresponding to $p(z_3, x_{1:3})$

Inference

$$p(x) = \alpha_1^\top \Lambda_2 \cdots \Lambda_T \mathbf{1}$$



- ▶ Each node corresponds to a state
- ▶ Each edge to an entry in the transition operator matrix

Scaling HMMs

Lessons from Large Neural Language Models

Large models perform better but are . . .

1. Slow to train
2. Prone to overfitting

We must overcome these issues when scaling HMMs

3 Techniques for Training Large HMMs

- ▶ Compact neural parameterization

↑ Generalization

- ▶ State dropout

↑ Speed ↑ Generalization

- ▶ Block-sparse emission constraints

↑ Speed

- ▶ Will cover a fourth in the second part of this talk

Technique 1: Neural Parameterization

- ▶ The transition and emission matrices have Z^2 and ZX entries
- ▶ Causes the number of parameters to explode as the state size increases
- ▶ We instead use a low-dimensional decomposition of all conditional distributions that greatly reduces the number of parameters

Neural Parameterization: Softmax

For both the transition and emission matrices, we use a softmax parameterization, which assumes a nonlinear D -dimensional decomposition

$$W \propto \exp \left(U \times V^T \right)$$

with embeddings $U \in \mathbb{R}^{Z \times D}$, $V \in \mathbb{R}^{Z \times D}$ or $\mathbb{R}^{X \times D}$

- Can further parameterize U or $V = \text{MLP}(E_u)$

Technique 2: State Dropout

- ▶ Dropout is a common technique for regularizing neural networks
 - ▶ Reduces a network's reliance on a particular neuron
- ▶ Extend dropout to the states of an HMM
 - ▶ Encourage broad utilization of all states

State Dropout

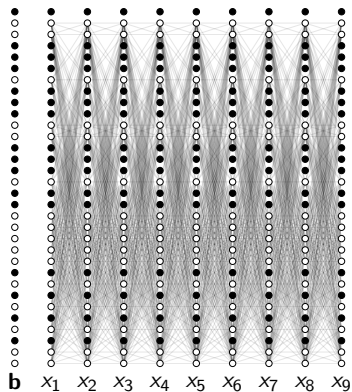
- ▶ At each batch, sample dropout mask $\mathbf{b} \in \{0, 1\}^Z$
- ▶ Compute distributional parameters by indexing into embeddings U, V

$$\left(\mathbf{b} \circ \begin{array}{|c|} \hline U_{\text{trans}} \\ \hline \end{array} \right) \times \left(\mathbf{b} \circ \begin{array}{|c|} \hline V_{\text{trans}} \\ \hline \end{array} \right)^{\top} \quad \left(\mathbf{b} \circ \begin{array}{|c|} \hline U_{\text{emit}} \\ \hline \end{array} \right) \times \begin{array}{|c|} \hline V_{\text{emit}}^{\top} \\ \hline \end{array}$$

(a) Unnormalized transition logits

(b) Unnormalized emission logits

State Dropout: Inference



- ▶ Shaded nodes depict dropped states
- ▶ Ignore dropped states during inference

Technique 3: Block-Sparse Emission Constraints

- ▶ Reduce cost of marginalization by enforcing structure
- ▶ Only allow each word to be emit by a subset of states
- ▶ Cost of inference is quadratic in the size of the largest subset due to sparsity

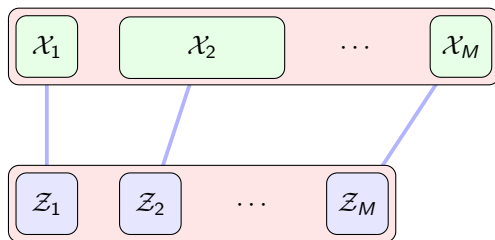
Block-Sparse Emission Constraints: Alignment

Start with a joint partitioning of both states and words

Indices $m \in [M]$

State partitions \mathcal{Z}_m

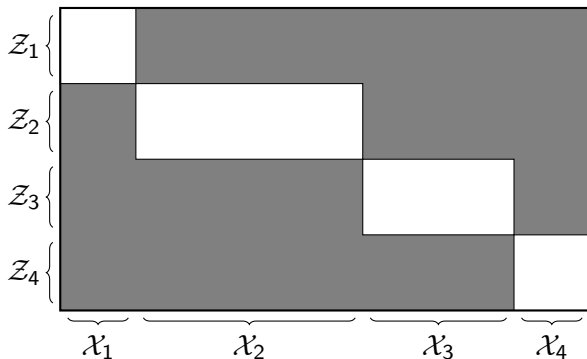
Word partitions \mathcal{X}_m



Block-Sparse Emission Constraints

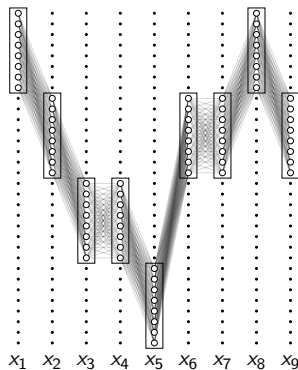
Given the unnormalized emission logits,

- ▶ Mask out unaligned state-word entries
- ▶ Normalize rows across words in aligned partition

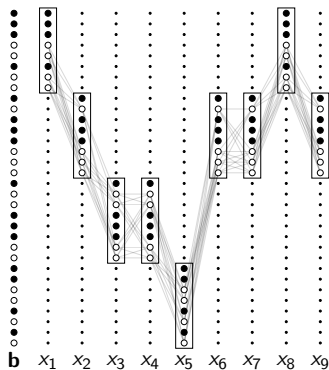


Block-Sparse Emissions: Inference

Given each word x_t , only the states in the correct group can occur



(a) Block-sparse emission



(b) With state dropout

Method Recap

- ▶ Compact neural parameterization

↑ Generalization

- ▶ State dropout

↑ Speed ↑ Generalization

- ▶ Block-sparse emission constraints

↑ Speed

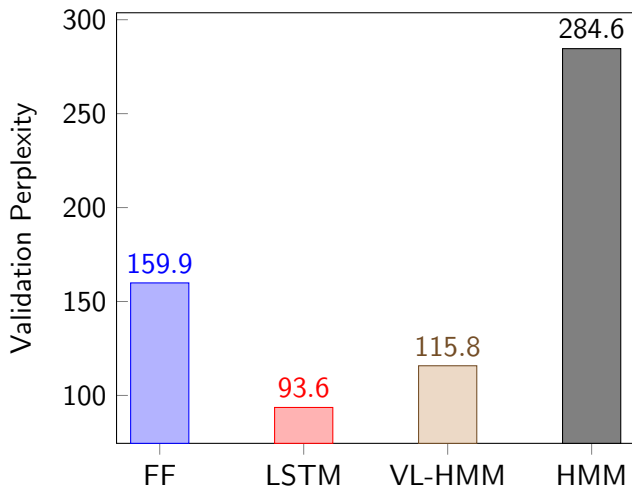
- ▶ A fourth after experiments

Experiments


Experiments

- ▶ Language modeling on Penn Treebank
- ▶ Baselines
 - ▶ Feedforward 5-gram model
 - ▶ 2-layer LSTM
 - ▶ A 900 state HMM (Buys et al 2018)
- ▶ Model
 - ▶ 2^{15} (32k) state very large HMM (VL-HMM)
 - ▶ $M = 128$ groups (256 states per type), obtained via Brown Clustering
 - ▶ Dropout rate of 0.5 during training

Results on PTB Validation Data



State Size Ablation



Validation perplexity on PTB by state size ($\lambda = 0.5$ and $M = 128$)

Other Ablations

| Model | Param | Train | Val |
|---------------------|-------|-------|-----|
| VL-HMM (2^{14}) | 7.2M | 115 | 134 |
| - neural param | 423M | 119 | 169 |
| - state dropout | 7.2M | 88 | 157 |

Discussion

- ▶ Greatly scaled the state size of HMMs
- ▶ Performance improved with increasing state size
- ▶ Still a large gap between RNNs and HMMs
- ▶ Does the emission sparsity constraint improve computation complexity at the price of accuracy?

Speeding up HMMs with Low-Rank Decompositions

Fast Inference with Low-Rank Decompositions

- ▶ The previous approach relied a pre-specified emission sparsity constraint
- ▶ Can we scale inference with a weaker constraint?
- ▶ Exploit structure in the transition matrix to speed up inference

Inference

Start by unpacking inference to reveal the most expensive step

$$p(x) = \alpha_1^\top \Lambda_2 \Lambda_3 \cdots \Lambda_T \mathbf{1}$$

with

$$\begin{aligned} \text{start,} \quad & [\alpha_1]_{z_1} = p(x_1 \mid z_1)p(z_1), \\ \text{and transition operators,} \quad & [\Lambda_t]_{z_{t-1}, z_t} = p(x_t \mid z_t)p(z_t \mid z_{t-1}) \end{aligned}$$

Inference

Decompose transition operators into transition matrix A and emission matrix O

$$\begin{aligned} p(x) &= \alpha_1^\top \Lambda_2 \cdot \Lambda_T \mathbf{1} \\ &= \alpha_1^\top (A \operatorname{diag}([O]_{\cdot, x_2})) \cdots \Lambda_T \mathbf{1} \\ &= \alpha_1^\top A \operatorname{diag}([O]_{\cdot, x_2}) \cdots A \operatorname{diag}([O]_{\cdot, x_T}) \mathbf{1} \end{aligned}$$

where the most expensive steps are the matrix-vector products $\alpha_t^\top A$, which take $O(Z^2)$ computation

Fast Matrix-Vector Products

- ▶ Goal is to reduce the naive matvec complexity of $O(Z^2)$
- ▶ Various methods
 - ▶ Sparsity (nnz entries)
 - ▶ Fast Fourier Transform ($Z \log Z$)
 - ▶ Low-Rank decomposition (ZR)
- ▶ We utilize low-rank decompositions

Low-Rank Factorization

Factor transition matrix A into product of skinny matrices

$$\boxed{\alpha^\top} \times \boxed{A} = \boxed{\alpha^\top} \times \boxed{U} \times \boxed{V^\top}$$

resulting in two matrix-vector products of cost $O(ZR)$ each

- ▶ Constraint: Entries of A must be nonnegative
- ▶ Solution: Use a nonnegative matrix factorization (NMF)

$$A = \phi(U)\phi(V)^\top,$$

with $\phi : \mathbb{R}^D \rightarrow \mathbb{R}_+$

Method Recap

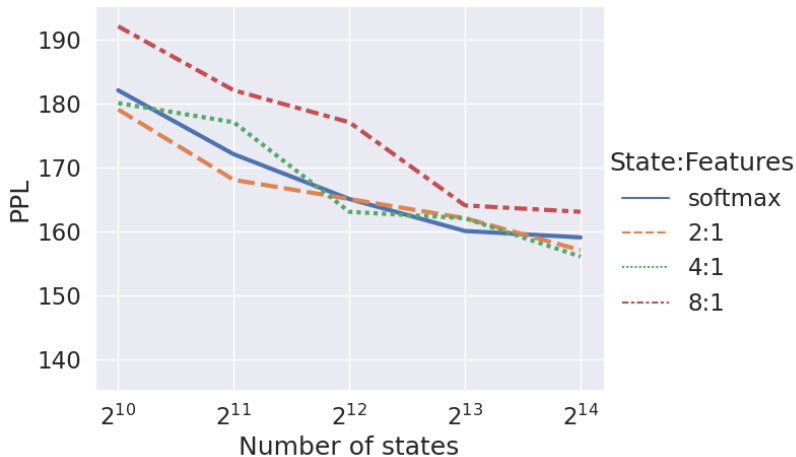
- ▶ Target key $O(Z^2)$ matvec step in inference
- ▶ Use NMF to reduce cost to $O(ZR)$
- ▶ How small can R be relative to Z without sacrificing accuracy?

Experiments

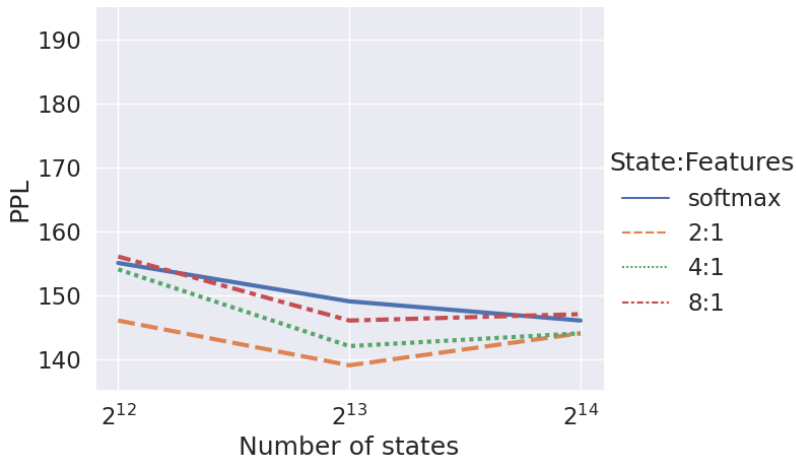
Experiments

- ▶ Language modeling on PTB
- ▶ Feature map $\phi(\mathbf{x}) = \exp(W\mathbf{x})$, with learned $W \in \mathbb{R}^{D \times R}$
- ▶ No sparsity constraints

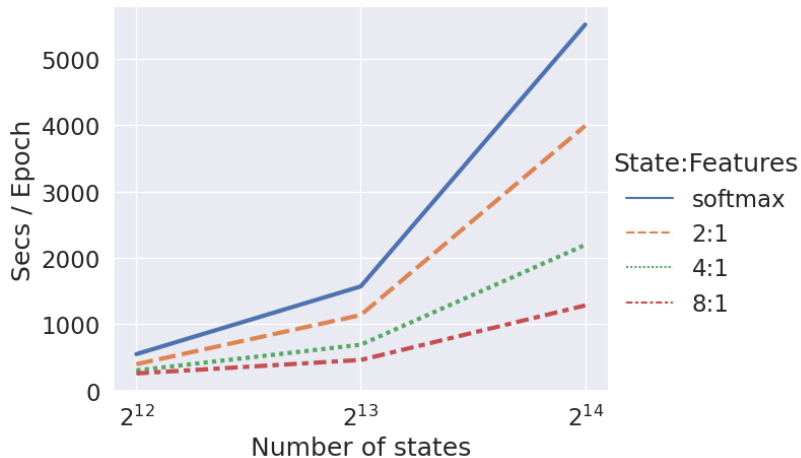
Scaling on PTB (Validation)



Further Scaling on PTB with Dropout (Validation)



Speed Comparison



Conclusion (TODO)

- ▶ Hopeful that HMMs can be competitive language models
- ▶ Introduced 4 techniques for tackling speed and overfitting
- ▶ Future work will extend to other discrete latent variable models

EOS

Citations

Generalized Softmax

- Softmax

$$p(z_t \mid z_{t-1}) = \frac{\exp(\mathbf{u}_{z_{t-1}}^\top \mathbf{v}_{z_t})}{\sum_z \exp(\mathbf{u}_{z_{t-1}}^\top \mathbf{v}_z)}$$

- Generalized Softmax

$$p(z_t \mid z_{t-1}) = \frac{K(\mathbf{u}, \mathbf{v})}{\sum_z K(\mathbf{u}, \mathbf{v}_z)} = \frac{\phi(\mathbf{u})^\top \phi(\mathbf{v})}{\sum_z \phi(\mathbf{u})^\top \phi(\mathbf{v}_z)},$$

for positive kernel $K : \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}_+$ and feature map $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^R$

Generalized Softmax: Inference

- The key $O(Z^2)$ step in the forward algorithm:

$$p(z_t \mid x_{<t}) = \sum_{z_{t-1}} p(z_t \mid z_{t-1}) p(z_{t-1} \mid x_{<t})$$

- In matrix form,

$$\gamma_t = \underbrace{\alpha_{t-1}}_{\mathbb{R}^Z} \underbrace{\Lambda}_{\mathbb{R}^{Z \times Z}},$$

where we have the probability of the

| | |
|-------------------------|--|
| current state, | $[\gamma_t]_{z_t} = p(z_t \mid x_{<t}),$ |
| last state, | $[\alpha_{t-1}]_{z_{t-1}} = p(z_{t-1} \mid x_{<t}),$ |
| transition probability, | $[\Lambda]_{z_{t-1}, z_t} = p(z_t \mid z_{t-1})$ |

Generalized Softmax: Inference

- Use generalized softmax in transition distribution

$$[\Lambda]_{z_{t-1}, z_t} = p(z_t \mid z_{t-1}) \propto \phi(\mathbf{u}_{z_{t-1}})^\top \phi(\mathbf{v}_{z_t})$$

- Allows us to apply associative property of matrix multiplication

$$\begin{aligned}\gamma_t &= \alpha_{t-1} \Lambda \\ &= \alpha_{t-1} (\text{diag}(d) \phi(U) \phi(V)^\top) \\ &= \underbrace{(\alpha_{t-1} \circ d)}_{\mathbb{R}^Z} \underbrace{\phi(U)}_{\mathbb{R}^{Z \times f}} \underbrace{\phi(V)^\top}_{\mathbb{R}^{f \times Z}},\end{aligned}$$

with stacked embeddings $\phi(U), \phi(V) = [\phi(\mathbf{v}_1), \dots, \phi(\mathbf{v}_Z)]$
and normalizing constants d

- Takes $O(Zf)$ time from left to right!