

# Scaling Hidden Markov Models

April 27, 2021

# Latent Variables Models in NLP

- ▶ NLP benchmarks are dominated by fully observed models
  - ▶ Transformers
  - ▶ Previously, Recurrent Neural Networks
- ▶ We instead explore latent variable models

# Latent Variable Models: Motivation

- ▶ LVMs posit a generative process involving unseen variables
- ▶ Maintain uncertainty over latent representations, rather than just output correlations
- ▶ Often improves interpretability and controllability
- ▶ Bottlenecked by the computational complexity of inference

## Research Question

To what extent is the performance of tractable latent variable models limited by scale and choices in parameterization?

# Language Modeling

To answer this question, we focus on the task of language modeling

- ▶ Given the words seen so far, predict the next word
- ▶ Language requires modeling long-range phenomena

How now, brown \_\_\_\_\_

# Hidden Markov Models in NLP

- ▶ Simplest latent variable models for time series data
- ▶ Are thought to be very poor language models
- ▶ We show they are better than previously thought, once scaled

## Background: Hidden Markov Models


# Hidden Markov Models (HMMs)

- ▶ Classical models for unsupervised per-word tag induction in NLP
  - ▶ Part-of-speech induction
  - ▶ Word alignment for translation
- ▶ Generative model separates temporal dynamics from emissions
  - ▶ First picks a sequence of latent states
  - ▶ Emits words given only the corresponding state



# Hidden Markov Models (HMMs)

For times  $t$ , model states  $z_t \in [Z]$ , and tokens  $x_t \in [X]$ ,



This yields the joint distribution

$$p(x, z) = \prod_t p(x_t \mid z_t) p(z_t \mid z_{t-1})$$

with

start state	$p(z_1),$
transitions	$p(z_t \mid z_{t-1}),$
and emissions	$p(x_t \mid z_t)$

# Inference

Given observed  $x = (x_1, \dots, x_T)$  We wish to maximize

$$p(x) = \sum_{z_1} \cdots \sum_{z_T} p(x, z) = \alpha_1^\top \Lambda_2 \Lambda_3 \cdots \Lambda_T \mathbf{1},$$

where we have the

$$\begin{aligned} \text{start,} \quad & [\alpha_1]_{z_1} = p(x_1 \mid z_1)p(z_1), \\ \text{and transition operators,} \quad & [\Lambda_t]_{z_{t-1}, z_t} = p(x_t \mid z_t)p(z_t \mid z_{t-1}) \end{aligned}$$

The result of each matvec has the alphas of the forward algorithm, i.e.  $\alpha_3 = \alpha_1 \Lambda_2 \Lambda_3$  has entries corresponding to  $p(z_3, x_{1:3})$

# Transition Operators

skip slide? mainly to make the connection between distributional parameters and inference really concrete.

- ▶ The transition operators  $\alpha_1, \Lambda_t$  are functions of the distributional parameters: starting state  $\pi$ , state transitions  $A$ , and emissions  $O$
- ▶ In particular

the start


$$\alpha_1 = \pi \circ [O]_{\cdot, x_1},$$

and the transition operators

$$\Lambda_t = A \text{diag}([O]_{\cdot, x_t})$$

# Inference

Given the start and transition operators, we can visualize inference



- ▶ Each node corresponds to a state
- ▶ Each edge to an entry in the transition operator matrix

## Scaling HMMs

# Lessons from Large Neural Language Models

Large models perform better but are . . .

1. Slow to train
2. Prone to overfitting

We must overcome these issues when scaling HMMs

### 3 Techniques for Training Large HMMs

- ▶ Compact neural parameterization

↑ Generalization

- ▶ State dropout

↑ Speed    ↑ Generalization

- ▶ Block-sparse emission constraints

↑ Speed

- ▶ Will cover a fourth in the second part of this talk

# Technique 1: Neural Parameterization

- ▶ The transition  $A$  and emission  $O$  matrices have  $Z^2$  and  $ZX$  entries
- ▶ Causes the number of parameters to explode as the state size increases
- ▶ We instead use a low-dimensional factorization of all conditional distributions that greatly reduces the number of parameters



# Neural Parameterization

For both the transition and emission matrices, we use a nonlinear  $D$ -dimensional decomposition

$$W \propto \exp \left( U \times V^T \right)$$

with embeddings  $U \in \mathbb{R}^{Z \times D}$ ,  $V \in \mathbb{R}^{Z \times D}$  or  $\mathbb{R}^{X \times D}$

- Can further parameterize  $U$  or  $V = \text{MLP}(E_u)$

## Technique 2: State Dropout

- ▶ Dropout is a common technique for regularizing neural networks
  - ▶ Reduces a network's reliance on a particular neuron
- ▶ Extend dropout to the states of an HMM
  - ▶ Encourage broad utilization of all states

# State Dropout

- ▶ At each batch, sample dropout mask  $\mathbf{b} \in \{0, 1\}^Z$
- ▶ Compute distributional parameters by indexing into embeddings  $U, V$


$$\left( \mathbf{b} \circ \begin{array}{|c|} \hline U_A \\ \hline \end{array} \right) \times \left( \mathbf{b} \circ \begin{array}{|c|} \hline V_A \\ \hline \end{array} \right)^T \quad \left( \mathbf{b} \circ \begin{array}{|c|} \hline U_O \\ \hline \end{array} \right) \times \begin{array}{|c|} \hline V_O \\ \hline \end{array}^T$$

(a) Unnormalized transition logits

(b) Unnormalized emission logits

# State Dropout: Inference

The cost of inference is reduced by state dropout




- ▶ Shaded nodes depict dropped states
- ▶ Ignore dropped states during inference

## Technique 3: Block-Sparse Emission Constraints

- ▶ Reduce cost of marginalization by enforcing structure
- ▶ Only allow each word to be emit by a subset of states
- ▶ Obtain an alignment by partitioning words and states jointly

# Block-Sparse Emission Constraints: Inference

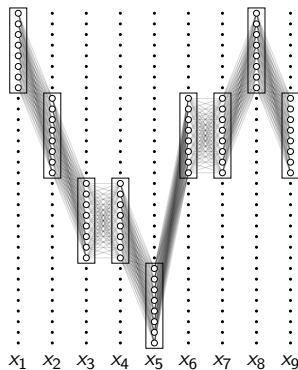
- ▶ Indices  $m \in [M]$
- ▶ State partitions  $\mathcal{Z}_m$
- ▶ Word partitions  $\mathcal{X}_m$



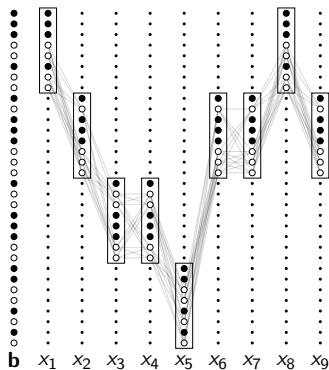


# Block-Sparse Emissions: Inference

Given each word  $x_t$ , only the states in the correct group can occur



(a) Block-sparse emission



(b) With state dropout



## Technique 4: Generalized Softmax

Focusing on the transition distribution,

- Softmax

$$p(z_t \mid z_{t-1}) = \frac{\exp(\mathbf{u}_{z_{t-1}}^\top \mathbf{v}_{z_t})}{\sum_z \exp(\mathbf{u}_{z_{t-1}}^\top \mathbf{v}_z)}$$

- Generalized Softmax

$$p(z_t \mid z_{t-1}) = \frac{K(\mathbf{u}, \mathbf{v})}{\sum_z K(\mathbf{u}, \mathbf{v}_z)} = \frac{\phi(\mathbf{u})^\top \phi(\mathbf{v})}{\sum_z \phi(\mathbf{u})^\top \phi(\mathbf{v}_z)},$$

for positive kernel  $K : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}_+$  and feature map  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^f$

# Generalized Softmax: Inference

- ▶ The key  $O(Z^2)$  step in the forward algorithm:

$$p(z_t \mid x_{<t}) = \sum_{z_{t-1}} p(z_t \mid z_{t-1}) p(z_{t-1} \mid x_{<t})$$

- ▶ In matrix form,

$$\gamma_t = \underbrace{\alpha_{t-1}}_{\mathbb{R}^Z} \underbrace{\Lambda}_{\mathbb{R}^{Z \times Z}},$$

where we have the probability of the

current state,	$[\gamma_t]_{z_t} = p(z_t \mid x_{<t}),$
last state,	$[\alpha_{t-1}]_{z_{t-1}} = p(z_{t-1} \mid x_{<t}),$
transition probability,	$[\Lambda]_{z_{t-1}, z_t} = p(z_t \mid z_{t-1})$

# Generalized Softmax: Inference

- ▶ Use generalized softmax in transition distribution

$$[\Lambda]_{z_{t-1}, z_t} = p(z_t \mid z_{t-1}) \propto \phi(\mathbf{u}_{z_{t-1}})^\top \phi(\mathbf{v}_{z_t})$$

- ▶ Allows us to apply associative property of matrix multiplication

$$\begin{aligned}\gamma_t &= \alpha_{t-1} \Lambda \\ &= \alpha_{t-1} (\text{diag}(d) \phi(U) \phi(V)^\top) \\ &= \underbrace{(\alpha_{t-1} \circ d)}_{\mathbb{R}^Z} \underbrace{\phi(U)}_{\mathbb{R}^{Z \times f}} \underbrace{\phi(V)^\top}_{\mathbb{R}^{f \times Z}},\end{aligned}$$

with stacked embeddings  $\phi(U), \phi(V) = [\phi(\mathbf{v}_1), \dots, \phi(\mathbf{v}_Z)]$   
and normalizing constants  $d$

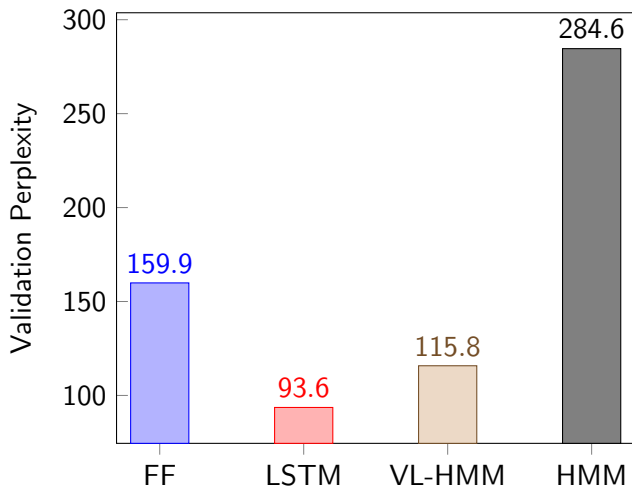
- ▶ Takes  $O(Zf)$  time from left to right!

## Experiments


# Experiments

- ▶ Language modeling on Penn Treebank
- ▶ Baselines
  - ▶ Feedforward 5-gram model
  - ▶ 2-layer LSTM
  - ▶ A 900 state HMM (Buys et al 2018)
- ▶ Model
  - ▶  $2^{15}$  (32k) state very large HMM (VL-HMM)
  - ▶  $M = 128$  groups (256 states per type), obtained via Brown Clustering
  - ▶ Dropout rate of 0.5 during training

## Results on PTB Validation Data



# State Size Ablation



Validation perplexity on PTB by state size ( $\lambda = 0.5$  and  $M = 128$ )

## Other Ablations

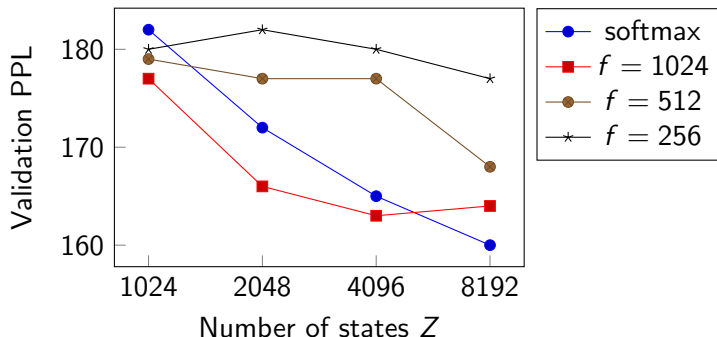
Model	Param	Train	Val
VL-HMM ( $2^{14}$ )	7.2M	115	134
- neural param	423M	119	169
- state dropout	7.2M	88	157



# Experiments

- ▶ Language modeling on PTB
- ▶ Work directly with feature map  $\phi(\mathbf{x}) = \exp\left(W\mathbf{x} - \frac{\|\mathbf{x}\|^2}{2}\right)$ ,  
with learned  $W \in \mathbb{R}^{d \times f}$
- ▶ No dropout or sparsity constraints

# Results on PTB Validation



- ▶ Holding number of features fixed, perplexity mostly improves or remains the same with an increasing number of states
- ▶ Achieve similar performance as softmax with around 4:1 state to feature ratio (also holds for 8k and 16k states)

# Conclusion

- ▶ Hopeful that HMMs can be competitive language models
- ▶ Introduced 4 techniques for tackling speed and overfitting
- ▶ Future work will extend to other discrete latent variable models

EOS

# Citations